

Descriptive Statistics in R

Victor A. Ordu

2021-07-16

Introduction

- ▶ Most of the functions related to pure statistics can be found in the *base* and *stats* package
- ▶ They are always pre-loaded by default in every R session

```
search()
```

```
## [1] ".GlobalEnv"          "package:stats"        "package:graphics"
## [4] "package:grDevices"    "package:utils"        "package:datasets"
## [7] "package:methods"      "Autoloads"            "package:base"
```

Summarization of Data

Summation

► Numerical data

The sum:

$$Sum(X) = \sum_{i=1}^N x_i$$

where x_i is a single element in a numeric vector and N is the number of elements in that vector.

We add up the values with `sum()`:

```
x <- 1:100  
sum(x)
```

```
## [1] 5050
```

► Categorical data

Here we do **counting** with `length()`

How many letters are there in the English alphabet?

```
length(LETTERS)
```

```
## [1] 26
```

Scaling

- ▶ Addition/subtraction
- ▶ Multiplication and/or division.
- ▶ When you apply a constant number, it computes on every element of the vector

```
mix <- c(5, 0, 28, -5, 29, 77, 10, 57, 28, 88, 298)
```

```
mix + 5
```

```
## [1] 10 5 33 0 34 82 15 62 33 93 303
```

```
mix * 2
```

```
## [1] 10 0 56 -10 58 154 20 114 56 176 596
```

```
mix / 3
```

```
## [1] 1.666667 0.000000 9.333333 -1.666667 9.666667 2
```

```
## [8] 19.000000 9.333333 29.333333 99.333333
```

```
mix %% 2
```

Maxima and minima

- ▶ Get the highest or lowest value in a collection.
- ▶ Functions: `max()` and `min()`

```
max(mix)
```

```
## [1] 298
```

```
min(mix)
```

```
## [1] -5
```


Summarization

- ▶ Tukey's five-number summary: minimum, maximum, mean, 1st quartile, 3rd quartile
- ▶ At a glance, gives you an idea of central statistic, dispersion and extremes of the data.

```
fivenum(mix)
```

```
## [1] -5.0  7.5 28.0 67.0 298.0
```

Sorting

Create an ordered array of values

```
sort(mix)
```

```
## [1] -5  0  5 10 28 28 29 57 77 88 298
```

In descending order

```
sort(mix, decreasing = TRUE)
```

```
## [1] 298 88 77 57 29 28 28 10  5  0 -5
```

Ranking

Get the rank of values within a collection

```
order(mix)
```

```
## [1] 4 2 1 7 3 9 5 8 6 10 11
```

Compare with original vector for clarity:

```
comp <- data.frame(original = mix, ordered = order(mix))  
comp
```

```
##      original ordered  
## 1           5        4  
## 2           0        2  
## 3          28        1  
## 4          -5        7  
## 5          29        3  
## 6          77        9  
## 7          10        5  
## 8          57        8
```

Measures of Central Tendency

Arithmetic mean

Arithmetic mean, of a population μ is defined by the equation

$$\mu = \frac{\sum_{i=1}^N x_i}{N}$$

- Computed with the function `mean()`

```
mean(mix)
```

```
## [1] 55.90909
```

Median

The median, the *middle* value in an ordered array, X , with N numbers can be described by the formula

$$\text{median}(X) = x_{(N+1)/2}$$

and when considering even numbered sets, we have the formula

$$\text{median}(X) = \frac{x_{N/2} + x_{(N/2)-1}}{2}$$

—

► Computed with the function `median()`

```
median(mix)
```

```
## [1] 28
```

The Difference between mean and median

```
# Add an outlier to the vector 'mix'
```

```
mix[length(mix) + 1] <- 1000
```

```
mix
```

```
## [1] 5 0 28 -5 29 77 10 57 28 88
```

```
mean(mix)
```

```
## [1] 134.5833
```

```
median(mix)
```

```
## [1] 28.5
```

Mode

There is no function per se in R for computing the mode, but developing a custom function for this purpose is trivial and would be done on a case-by-case basis.

Measures of Dispersion

When you want to assess the variability in a set of data or a given variable

Range

The range, R , is the difference between the lowest value, x_S and the largest value, x_L , i.e.

$$R = x_L - x_S$$

In R, the function `range` returns both values as a 2-element numeric vector.

```
range(mix)
```

```
## [1] -5 1000
```

Interquartile Range

Is the range between the first and the third quartile - A measure of the “middle fifty percent” of a dataset - It is the range between the 75th and 25th percentile


$$IQR = Q_3 - Q_1$$

```
IQR(mix)
```

```
## [1] 71
```

In the box-and-whiskers plot, the “box” displays the *IQR*.

```
boxplot(mix)
```



Variance

Dispersion around the mean, with the difference squared to create absolute values.

$$\text{Var}(X) = \frac{\sum_{i=1}^N (x_i - \mu)^2}{N}$$

```
var(mix)
```

```
## [1] 80993.9
```

Standard Deviation

This is essentially (elementarily) the square root of the variance

$$s.d.(X) = \sqrt{\frac{\sum_{i=1}^N (x_i - \bar{x})^2}{N}}$$

```
sd(mix)
```

```
## [1] 284.5943
```

Coefficient of variation

Used to measure the variation between 2 different but related datasets It is the ratio between the standard deviation and the mean, expressed as a percentage.

$$C.V. = \frac{s.d.}{\bar{x}}(100)$$

Grouping of Data

Numerical vs. Categorical Data

- ▶ With numerical data we create intervals first, converting them to discrete categories.
- ▶ Then we create a frequency tabulation

```
cut(mix, 3) # create 3 groups using defaults
```

```
## [1] (-6,330]      (-6,330]      (-6,330]      (-6,330]      (-6,330]
## [7] (-6,330]      (-6,330]      (-6,330]      (-6,330]      (-6,330]
## Levels: (-6,330] (330,665] (665,1e+03]
```

```
cut(mix, c(-5, 0, 100, 1000)) # Note missing value
```

```
## [1] (0,100]      (-5,0]      (0,100]      <NA>      (0,100]
## [7] (0,100]      (0,100]      (0,100]      (0,100]      (10,100]
## Levels: (-5,0] (0,100] (100,1e+03]
```

,

```
cut(mix, c(-4, 0, 100, 1000), include.lowest = TRUE)
```

```
## [1] (0,100]      [-4,0]      (0,100]      <NA>      (0,100]
## [7] (0,100]      (0,100]      (0,100]      (0,100]      (10,100]
## Levels: [-4,0] (0,100] (100,1e+03]
```

```
cut(mix, breaks = c(min(mix), median(mix), max(mix)))
```

```
## [1] (-5,28.5]      (-5,28.5]      (-5,28.5]      <NA>  
## [6] (28.5,1e+03] (-5,28.5]      (28.5,1e+03] (-5,28.5]  
## [11] (28.5,1e+03] (28.5,1e+03]  
## Levels: (-5,28.5] (28.5,1e+03]
```

Frequency Distribution

- ▶ For categorical/discrete data
- ▶ In R, we use the `table()` function.

```
table(esoph$agegp)
```

```
##
```

```
## 25-34 35-44 45-54 55-64 65-74 75+
```

```
##    15    15    16    16    15    11
```

Cumulative Frequency Distribution

```
cumsum(mix)
```

```
##    [1]     5     5    33    28    57   134   144   201   229   317
```

```
# Compare with original
```

```
cumfrq <- data.frame(original = mix, cum.freq = cumsum(mix))  
cumfrq
```

##	original	cum.freq
## 1	5	5
## 2	0	5
## 3	28	33
## 4	-5	28
## 5	29	57
## 6	77	134
## 7	10	144
## 8	57	201
## 9	28	229
## 10	88	317
## 11	298	615
## 12	1000	1615

Something is wrong with this arrangement. Can you spot it?

Some important characteristics of vectors

- ▶ We need to discuss some important behaviours of vectors
- ▶ They have practical implications when you are carrying out analyses.

Vectorization

- ▶ A process whereby an operation carried out affects every element of a vector.
- ▶ First let's create a vector with random whole numbers

```
set.seed(134)
myvalues <- sample(1:7, size = 10, replace = TRUE)
myvalues
```

```
## [1] 4 2 2 3 6 4 2 3 2 7
```

Arithmetic

```
myvalues + 1
```

```
## [1] 5 3 3 4 7 5 3 4 3 8
```

```
myvalues * 5
```

```
## [1] 20 10 10 15 30 20 10 15 10 35
```

```
myvalues < 3
```

```
## [1] FALSE TRUE TRUE FALSE FALSE FALSE TRUE FALSE TR
```

Vectorized Logical 'AND'

```
set.seed(54)
newvalues <- sample(myvalues, size = length(myvalues), repl
identical(newvalues, myvalues)

## [1] FALSE
```

```
myvalues < newvalues
```

```
## [1] FALSE FALSE FALSE TRUE FALSE FALSE TRUE FALSE FAI
```

```
myvalues > 2 && myvalues < 5
```

```
## [1] TRUE
```

But with “Vectorized” & we do this check element by element.

```
myvalues > 2 & myvalues < 5
```

```
## [1] TRUE FALSE FALSE TRUE FALSE TRUE FALSE TRUE FA
```


- ▶ `&&` and `||` are best for logical single outcome computations
e.g. `is.character(x) && is.atomic(x)`
- ▶ `&` and `|` are used to compare elements one by one along the
length of two vectors. e.g. `is.na(x) | duplicated(x)`

Recycling

When 2 vectors are used together as operands in a computation e.g. $A + B$, if $\text{length}(A) \neq \text{length}(B)$, the shorter one will be recycled until the computation on the longer vector is completed, without prejudice to which element of the shorter is used last.

Thus,

```
x <- 5  
y <- 1:8  
x * y      # Recycling of length 1L akin to vectorization  
  
## [1]  5 10 15 20 25 30 35 40
```

```
x <- c(5, 10)
```

```
x * y
```

```
## [1]  5 20 15 40 25 60 35 80
```

Display side-by-side

```
data.frame(y, prod = x * y)
```

```
##    y prod
## 1 1     5
## 2 2    20
## 3 3    15
## 4 4    40
## 5 5    25
## 6 6    60
## 7 7    35
## 8 8    80
```

Demonstrate abrupt stop to recycling

```
x[3] <- 15L
```

```
x * y
```

```
## [1]  5 20 45 20 50 90 35 80
```

Viewed in a data frame...

##	x	y	x.times.y
## 1	5	1	5
## 2	10	2	20
## 3	15	3	45
## 4	5	4	20
## 5	10	5	50
## 6	15	6	90
## 7	5	7	35
## 8	10	8	80

Missing values

Most computations, such as the ones above will return NA if a missing value is present.

```
mix[7] <- NA  
mix
```

```
## [1] 5 0 28 -5 29 77 NA 57 28 88
```

```
sum(mix)
```

```
## [1] NA
```


To fix this, remove NA from the computation. See the `sum`'s signature:

```
args(sum)
```

```
## function (... , na.rm = FALSE)
```

```
## NULL
```

Where `...` received objects to be summed up and `na.rm` tunes our ability to discount NA from the computation.

Thus,

```
sum(mix, na.rm = TRUE)
```

```
## [1] 1605
```

Many other functions used for statistical computations have arguments for dealing with NA. Whether you leave or remove depends entirely on the circumstance at any given time.

```
mean(mix)
```

```
## [1] NA
```

```
mean(mix, na.rm = TRUE)
```

```
## [1] 145.9091
```

The real point is to always glance at the help files e.g. `?mean`. The function `args()` used earlier also provides a way to quickly look at a function signature, but it is not always guaranteed to be informative, especially with generic functions.

```
args(plot)
```

```
?plot
```